

Practical 2 – MATLAB Parallel Computing Toolkit

Cammaran Clark and Kian Frassek

EEE4021 Class of 20

University of Cape Town

South Africa

[†]CLRCAM007 [‡]FRSKIA001

Abstract—

I. INTRODUCTION

This report details an investigation into MATLAB's Parallel Computing Toolkit

A. Part 1 Question A

No, the `spmd` block does not necessarily parallelize code in the GPU. The `spmd` block or Run Single Programs on Multiple Data Sets is used to parallelize code across multiple workers in a parallel pool but not specifically the GPU. To use a GPU to process code you would use the matlab functions `gpuArray()` function `gatherfun()` will force the code to use the GPU.

B. Part 1 Question B

MISSIIIInG

II. METHODOLOGY OF PART 2

A. Question A

Create a bubble sort function that can sort an array of randomised values. The function will be used to sort the columns of a 100x100 matrix.

B. Question B

Using the bubble sort function created in the previous question, test it against the inbuilt MATLABsort function. The testing will be conducted on square matrices of sizes 100, 200, 500, 1000, 10000 and calculate the speed up from the user created function to MATLAB's inbuilt function. It is expected that the inbuilt function should be faster for all sizes, and the speed up should increase as the size goes up.

C. Question C

Using MATLABparallelism to create a new bubble sort function which using explicit parallelism using a 'parfor' or 'spmd' block The testing will be conducted on square matrices of sizes 100 and 5000 and calculate the speed up from the user created function to MATLAB's inbuilt function It is expected that the inbuilt function should be faster for all sizes, but the explicit parallelism bubble sort should be faster than the user created bubble sort without parallelism.

III. RESULTS OF PART 2

A. Question A

The bubble sort function was created based off of sudo code for a generalised bubble sort. The code is shown below

```
function array = BubbleSort(array)
for i = 1:length(array)
for j = length(array):-1:i+1
if array(j) > array(j-1)
temp = array(j);
array(j) = array(j-1);
array(j-1) = temp;
end
end
end
return;
end
```

This code took 0.0087 seconds to sort the 100x100 matrix

B. Question B

The code to test the user function is shown below. 'timetaken' is a matrix for storing the times of both the user and inbuilt functions.

```
% Initialize the cell array
testsizes = {100; 200; 500; 1000; 1200};
timetaken = cell(1+length(testsizes),4);
timetaken(1,:) = {'size', 'bubble sort time', 'inbuilt sort time', 'speed up'};
timetaken(2:end,1) = num2cell(testsizes); % should use [100; 200; 500; 1000; 10000] but takes too long so

% Loop over each row in timetaken
for i = 2:size(timetaken,1)
% Call the timetest function with the value in the first column of the i-th row
[timeBubble, timeInbuilt] = timetest(timetaken(i,1));
timetaken(i,2) = timeBubble;
timetaken(i,3) = timeInbuilt;
% Calculate speed up
timetaken(i,4) = timeBubble/timeInbuilt;
end
% Display the results
disp(timetaken)
```

The function 'timetest' is given below. It takes a matrix size, creates a square matrix, and then outputs the sort time for both the user and inbuilt functions.

```
function [t_user,t_inbuilt] = timetest(size)

X=rand(size, size);
Y=X;

tic
for i = 1:size
X(:,i) = BubbleSort(X(:,i));
end
t_user = toc();

tic
for i = 1:size
Y(:,i) = sort(Y(:,i));
end
t_inbuilt = toc();

display('Time taken by the bubble sort function was ' + t_user + ' s. Time taken by the inbuilt sort function was ' + t_inbuilt + ' s');

return;
end
```

The results from the tests are shown in the table below

TABLE I
USER VS INBUILT SORT FUNCTION

Matrix size	User	Inbuilt	Speed up
100	0.0087	0.0002	40.507
200	0.0205	0.0003	62.791
500	0.1780	0.0012	145.18
1000	2.1324	0.0048	445.16
1000	1813.9	1.5631	1160.4




Fig. 1: User vs inbuilt sort function

WHYYYYY

C. Question C

The code to test the parallelism user function is shown below. 'timetaken' is a matrix for storing the times of both the user and inbuilt functions.

```
testsizes = [100, 5000];
timetaken_par = cell(1+length(testsizes),4);
timetaken_par(1,:) = {"size", "bubble sort time", "inbuilt sort time", "speed up"};
timetaken_par(2:end,1) = num2cell(testsizes); % should use [100; 200; 500; 1000; 10000] but takes too long so using 1100 for know
% Loop over each row in timetaken_par
for i = 2:size(timetaken_par,1)
    % Call the timetest function with the value in the first column of the i-th row
    timeBubble = timebubble_parallelism(timetaken_par(i,1));
    timeInbuilt = timesort_parallelism(timetaken_par(i,1));
    timetaken_par(i,2) = timeBubble;
    timetaken_par(i,3) = timeInbuilt;
    % Calculate speed up
    timetaken_par(i,4) = timeBubble/timeInbuilt;
end
% Display the results
disp(timetaken_par)
```

The function 'timesort_parallelism' simply returns the time to sort a square matrix of a given size using the inbuilt sort function. The function 'timebubble_parallelism' is given below. It takes a matrix size, creates a square matrix, and then outputs the sort time for the user parallelism bubble-sort function.

```
function t_user = timebubble_parallelism(size)
X=rand(size, size);

tic
spmd
myStart = (spmdIndex - 1) * 25 + 1;
myEnd = myStart + 24;
for i = myStart:myEnd
    X(:,i) = BubbleSort(X(:,i));
end
end
t_user = toc();

display("Time tacken by the parallelism bubble sort function was " + t_user )

return;
end
```

The results from the tests are shown in the table below

TABLE II
USER VS INBUILT SORT FUNCTION

Matrix size	User	Inbuilt	Speed up
100	0.7921	0.0011	690.97
5000	0.2.1446	0.4808	4.4607




Fig. 2: User parallelism vs inbuilt sort function

WHYYYYY

IV. CONCLUSION

The conclusion should provide a summary of your findings. Many people only read the introduction and conclusion of a paper. They sometimes scan the tables and figures. If the conclusion hints at interesting findings, only then will they bother to read the whole paper.

You can also include work that you intend to do in future, such as ideas for further improvements, or to make the solution more accessible to the general user-base, etc.

Publishers often charge "overlength article charges" [?], so keep within the page limit. In EEE4084F we will simulate overlength fees by means of a mark reduction at 10% per page. Late submissions will be charged at 10% per day, or part thereof.